

Fine-Tuning GR00T N1.6 for Cross-Embodiment Adaptation: Ablating Backbone Importance in Low-Data Regimes for Small-Arm Manipulation

Chan, Tobias Yin-Ching

Chinese University of Hong Kong

Abstract

Foundation models like GR00T N1.6 represent a paradigm shift in robotics, enabling cross-embodiment generalization across diverse robot morphologies. However, adapting such models to low-cost, small-scale arms like the SO-ARM101 (SO101) under limited data and compute constraints remains underexplored. This proposal outlines a reproducible, simulation-first study that fine-tunes GR00T N1.6 on a dual-camera (front + wrist) pick-and-place task ("pick 3 oranges") using high-quality human teleoperation demos augmented with synthetic trajectories via Isaac Lab Mimic. We focus on a minimal ablation set to isolate the importance of fine-tuning different backbones: the Diffusion Transformer (DiT) action head, top Vision-Language Model (VLM) layers, and parameter-efficient adaptations like LoRA. Additionally, we incorporate camera ablations and a cross-embodiment initializer comparison (base GR00T vs. single-arm finetuned checkpoints) to test adaptation efficiency. Evaluations emphasize long-horizon success rates, robustness to perturbations, and VRAM feasibility on a single 24GB GPU. This work addresses a key gap: how humanoid-scale VLAs adapt to accessible hardware with small data, providing insights into selective fine-tuning strategies that balance performance, stability, and generality.

Introduction

Recent advances in vision-language-action (VLA) models, such as NVIDIA's GR00T series, have demonstrated strong zero-shot capabilities in humanoid robotics. GR00T N1.6, built on the Cosmos-Reason-2B VLM and a DiT action generator, supports cross-embodiment fine-tuning for new robot morphologies, making it a promising foundation for democratizing advanced manipulation. However, public workflows predominantly target larger platforms (e.g., Bridge or WidowX arms), leaving open questions about adaptation to low-cost embodiments like the SO-ARM101—a 6-DoF arm with dual RGB cameras (front for global context, wrist for precise grasping).

This study targets a practical cross-embodiment gap: "How does a humanoid-scale VLA adapt to a low-cost small arm with dual cameras under small data?" We hypothesize that, for in-distribution simulation tasks like repeated pick-and-place, fine-tuning the DiT action head is critical for embodiment mapping and sequencing, while partial VLM adaptation (top layers) enhances fusion of dual-camera inputs without necessitating full backbone updates—especially under 24GB VRAM constraints. To test this, we integrate a simulation-first pipeline (teleop demos + Mimic augmentation) with targeted ablations on backbone components, camera views, and initializers. Our approach emphasizes reproducibility, leveraging pinned versions of Isaac Sim/Lab and LeIsaac workflows, to produce "research-grade" results that inform efficient fine-tuning for resource-limited settings.

Motivation stems from GR00T's own evolution: N1 post-training freezes the VL language component to preserve generality; N1.5 freezes the entire VLM; N1.6 partially unfreezes top VLM layers during pretraining/finetuning. These choices suggest selective adaptation is not a compromise but a deliberate strategy to mitigate overfitting and catastrophic forgetting in low-data regimes. By ablating these in a dual-cam, small-arm context, we provide empirical evidence on "which backbone matters" for cross-embodiment transfer.

Related Work

GR00T N1.6 builds on prior VLAs like RT-1/2 and Octo, emphasizing diffusion-based action generation for smoother, long-horizon policies. Its post-training selectively trains DiT + top VLM layers (e.g., Cosmos-Reason-2B variant) while freezing lower layers to maintain visual/language priors. Public fine-tuning guides for SO10x arms highlight VRAM bottlenecks (~25GB default) and recommend disabling DiT tuning for <24GB setups, underscoring the need for efficient strategies.

Isaac Lab Mimic extends datasets by decomposing demos into subtasks (e.g., reach, grasp, lift) and generating variations via transformation/stitching, with success filtering. LeIsaac integrates SO101 teleop with Isaac Lab for HDF5 recording, supporting LeRobot v2 conversion for GR00T compatibility.

No public study isolates backbone importance for SO101 dual-cam on multi-step tasks, but GR00T papers motivate partial unfreezing: tuning vision is compute-intensive (batch size drops to ~16 on A6000) yet valuable for domain shifts. Our work fills this by combining Mimic-augmented data with ablations, testing if dual-cam observability reduces the need for deep visual adaptation.

Method

Pipeline Overview

We adopt a sim-first approach for coherence and feasibility, aligning with LeIsaac's SO101 → Isaac Lab workflow. The goal is to show that 50–200 human teleop episodes + Mimic-generated synthetics enable effective GR00T N1.6 fine-tuning as a "new embodiment," improving closed-loop success on harder-than-zero-shot tasks.

1. **Teleop Collection (SO101 Leader → Isaac Lab):** Collect ~50–200 successful episodes in a custom Isaac Lab task (e.g., LeIsaac-SO101-PickOrange-v0 extended to "pick 3 oranges"). Use `teleop_se3_agent.py --task=LeIsaac-SO101-PickOrange-v0 --teleop_device=so101leader --record --dataset_file=dataset.hdf5`. Each episode starts/ends at standardized poses, minimizes jerk, and includes language instructions (e.g., "Pick up three oranges and place them on the plate").
2. **Data Annotation & Augmentation:** Annotate subtasks (e.g., reach pregrasp, grasp, lift, move to target, release). Generate 1k–50k synthetics with Isaac Lab Mimic, ensuring Mimic-compatible env configs (e.g., boolean success checks).
3. **Dataset Conversion:** Convert HDF5 to LeRobot v2 format for GR00T compatibility.

4. **Fine-Tuning:** Train as NEW_EMBODIMENT with SO100/SO101 modality configs. Use selective strategies (detailed in Experiments) on a single RTX 4090 (24GB VRAM), with flags like `--no-tune_diffusion_model` for memory ablation.
5. **Evaluation:** Open-loop (action prediction error on held-out episodes); closed-loop (success rate in Isaac Lab under randomized conditions, e.g., object poses). Optional: DAgger in sim for compounding error fixes.

Data Collection

Each episode is a time-series with:

- **Essentials:** Dual RGB (front + wrist), joint positions/velocities (6-DoF arm + gripper), relative joint delta actions, timestamps, episode boundaries, language instruction.
- **Extras:** End-effector pose, gripper contact, depth, object poses, success labels.

Collect > needed, filter to clean successes via replay scripts. Log per-episode metadata (e.g., success, notes, seeds) in JSON for hygiene.

Task Selection

We select a Tier B/C task: "Pick 3 oranges and place on plate" (long-horizon, contact-rich, partial occlusions). This benefits from Mimic variations and tests sequencing/recovery, where dual-cam shines (front: layout; wrist: grasp detail). Baseline: single pick (Tier A).

Environment Setup

- **Isaac Lab/Sim:** Ubuntu, Isaac Sim 4.5.0 (pip), Isaac Lab v2.1.0 (pinned commit).
- **SO-ARM101:** Leader arm via serial/USB for teleop.
- **GR00T Codebase:** NVIDIA/Isaac-GR00T (pinned commit per Seed guide).

This ensures "reproduce in one day" via README.

Experiments

Minimal Ablation Set

Hold dataset fixed (~30 episodes human + Mimic-augmented to 1k total), training steps fixed, on single 24GB GPU. Compare:

1. **DiT-Only:** Freeze all VLM. Train only DiT + embodiment adapters/encoders/decoders. Tests: "Vision is good enough; policy needs motor learning."
2. **N1.6 Default-Style Partial VLM:** Tune top 4 VLM layers + DiT (as per DeepWiki), freeze lower VLM. Tests: "Do we need backbone adaptation for dual-cam + SO101?"
3. **LoRA on Top VLM + DiT:** Same targets as (2), but with LoRA (instead of full gradients). Isolates: "Benefit from extra capacity vs. dense weight updates?"

4. **Camera Ablations:** Evaluate (1–3) under dual-cam, front-only, wrist-only. Tests vision hunch: If "vision encoder matters less," expect small deltas; anticipate larger in robustness/recovery.

Cross-Embodiment Twist

Start from base nvidia/GR00T-N1.6-3B *and* a single-arm finetuned checkpoint (e.g., Bridge/WidowX from Isaac-GR00T repo). Compare sample efficiency adapting to SO101. Tests: "Is a single-arm finetuned model a better initializer for another arm than the base cross-embodiment model?"

Metrics

- **Long-Horizon Specific:** Probability of completing 1/2/3 oranges; failure modes (grasp vs. place vs. transition); sensitivity to perturbations (object positions).
- **Standard:** Closed-loop success rate (randomized sim conditions); open-loop action error.
- **Compute:** VRAM usage, training time, batch size.

Run 3 seeds per config for statistical significance.

Expected Results and Discussion

We expect DiT-only to suffice for basic embodiment mapping but underperform on sequencing/occlusions, where partial VLM + dual-cam boosts recovery (e.g., +20–30% on 3-orange completion). LoRA should approximate partial VLM at lower VRAM, validating it for <24GB setups. Cross-embodiment: Single-arm initializers may show faster convergence but risk embodiment-specific biases. If vision adaptation matters less in sim (in-distribution), wrist-only could rival dual-cam on grasps but falter on layout—highlighting fusion's role.

Limitations: Sim-only (real extension future); small data risks overfitting. This frames selective tuning as optimal for low-resource cross-embodiment, countering "full fine-tune everything" assumptions.

Conclusion

This proposal delivers a gap-filling study on GR00T N1.6 adaptation, combining efficient data pipelines with backbone ablations to guide practical robotics. By publishing code/datasets, it advances accessible, reproducible research.

References

1. GR00T N1: An Open Foundation Model for Generalist Humanoid Robots (arxiv.org/abs/2503.14734).
2. Fine-tune Isaac GR00T N1.5 for LeRobot SO-101 Arm and Deploy on Jetson Thor (wiki.secdstudio.com).
3. GR00T N1.5 (research.nvidia.com).

4. GR00T N1.6 (research.nvidia.com).
5. Fine-tuning GR00T (deepwiki.com/nvidia/isaac-gr00t/fine-tuning-gr00t).
6. Building Generalist Humanoid Capabilities with NVIDIA Isaac GR00T N1.6 Using a Sim-to-Real Workflow (developer.nvidia.com).
7. NVIDIA/Isaac-GR00T: NVIDIA Isaac GR00T N1.6 - A Foundation Model for Generalist Robots (github.com/NVIDIA/Isaac-GR00T).
8. Isaac Lab Mimic Documentation (isaac-sim.github.io/isaac_lab).
9. LeIsaac Workflow (wiki.seeedstudio.com).